

# Spark Project

April 24, 2018

## 1 Spark Assignment 2018

### 1.1 Exercise 1

For each exercise, the first cell creates the function and the second cell executes the query and prints the result.

#### 1.1.1 Exercise 1.1.

```
In [1]: from pyspark.sql import SQLContext
import pandas as pd
from pyspark.sql.functions import mean, min, max, col, avg, monotonically_increasing_id
import pyspark.sql.functions as F
from operator import add
import time

file = "pagecounts-20160101-000000"
sqlContext = SQLContext(sc)
rdd = sc.textFile(file)

def f1(rdd):
    input_data = rdd.map(lambda line: line.split(" "))
    df_normal = input_data.toDF()

    # appearance
    df = df_normal.limit(15).toPandas()
    for index, row in df.iterrows():
        print(row[0], row[1], row[2], row[3])

In [4]: f1(rdd)

aa %CE%92%CE%84_%CE%95%CF%80%CE%B9%CF%83%CF%84%CE%BF%CE%BB%CE%AE_%CE%99%CF%89%CE%AC%CE%BD%CE%B
aa %CE%98%CE%B5%CF%8C%CE%B4%CF%89%CF%81%CE%BF%CF%82_%CE%91%CE%84_%CE%9B%CE%AC%CF%83%CE%BA%CE%B
aa %CE%9C%CF%89%CE%AC%CE%BC%CE%B5%CE%B8_%CE%95%CE%84/e1/%CE%9C%CE%B5%CF%87%CE%BC%CE%AD%CF%84_%
aa %CE%A0%CE%B9%CE%B5%CF%81_%CE%9B%27_%CE%91%CE%BD%CF%86%CE%AC%CE%BD/e1/%CE%A0%CE%B9%CE%B5%CF%
aa %CE%A3%CE%A4%CE%84_%CE%A3%CF%84%CE%B1%CF%85%CF%81%CE%BF%CF%86%CE%BF%CF%81%CE%AF%CE%B1/e1/%C
aa %D0%A1%D0%BE%D0%BB%D0%B8_484_%D0%BF.%D0%BC 1 4750
aa 271_a.C 1 4675
```

```

aa Battaglia_di_Qade%C5%A1/it/Battaglia_dell%27Oronte 1 4765
aa Category:User_th 1 4770
aa Chiron_Elias_Krase 1 4694
aa County_Laois/en/Queen%27s_County,_Ireland 1 4752
aa Dassault_rafaele 2 9372
aa Dyskusja_wikiprojektu:Formu%C5%82a_1/%22/pl/Polacy_w_Formule_1%22 1 4824
aa E.Desv 1 4662
aa Enclos-apier/fr/Enclos-Apiers_en_C%C3%B4te_d%27Azur 1 4772

```

### 1.1.2 Exercise 1.2.

```

In [5]: def f2(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code','title','hits', 'size'])
        print(df_normal.count())

```

```

In [6]: f2(rdd)

```

5046226

### 1.1.3 Exercise 1.3.

```

In [2]: def f3(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code','title','hits', 'size'])

        # convert to double
        df_normal = df_normal.withColumn('size', col('size').cast('double'))
        df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

        # get stats
        x = df_normal.select([mean('size'), min('size'), max('size')]).collect()
        result = "Mean = " + str(x[0][0]) + "; Min = " + str(x[0][1]) + "; Max = " + str(x[0][2])
        return result

```

```

In [8]: print(f3(rdd))

```

Mean = 101423.92964801814; Min = 0.0; Max = 141180155987.0

### 1.1.4 Exercise 1.4.

```

In [9]: def f4(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code','title','hits', 'size'])

        # convert to double from string

```

```

df_normal = df_normal.withColumn('size', col('size').cast('double'))
df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

#get max
max_value = df_normal.select(max('size')).collect()[0][0]
l=[max_value]
df = df_normal.where(df_normal.size.isin(l))

df= df.toPandas()
for index, row in df.iterrows():
    print(row[0], row[1], row[2], row[3])

```

In [10]: f4(rdd)

en.mw en 5466346.0 141180155987.0

Only one record with maximum page size value.

### 1.1.5 Exercise 1.5.

```

In [3]: def f5(rdd):
    input_data = rdd.map(lambda line: line.split(" "))
    df_normal = input_data.toDF(['code','title','hits', 'size'])

    # convert to double from string
    df_normal = df_normal.withColumn('size', col('size').cast('double'))
    df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

    # get max
    max_value = df_normal.select(max('size')).collect()[0][0]
    l=[max_value]
    df = df_normal.where(df_normal.size.isin(l))

    # pick most hits if several records
    max_hits = df.select(max('hits')).collect()[0][0]
    l=[max_hits]
    df = df_normal.where(df.hits.isin(l))

    return df

```

```

In [12]: df = f5(rdd).toPandas()
for index, row in df.iterrows():
    print(row[0], row[1], row[2], row[3])

```

en.mw en 5466346.0 141180155987.0

### 1.1.6 Exercise 1.6.

```
In [4]: def f6(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

        # create title length column
        new = df_normal.withColumn("length", F.length('title'))

        # get longest page title
        max_value = new.select(max('length')).collect()[0][0]
        l=[max_value]
        df = new.where(new.length.isin(l))

        return df
```

Result of query is commented out because query is very long. Please remove "#" to execute query.

```
In [15]: #df = f6(rdd).toPandas()
        #for index, row in df.iterrows():
        #    print(row[1])
        #print()
```

### 1.1.7 Exercise 1.7.

```
In [5]: def f7(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

        # convert to double from string
        df_normal = df_normal.withColumn('size', col('size').cast('double'))
        df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

        # get mean
        mean_value = df_normal.select(mean('size')).collect()[0][0]

        # filter
        df = df_normal.filter(df_normal.size > mean_value)

        df = df.rdd
        return df
```

```
In [16]: f7(rdd).take(5)
```

```
Out[16]: [Row(code='aa', title='Main_Page', hits=5.0, size=266946.0),
          Row(code='ab', title='%D0%90%D0%B2%D0%B8%D0%BA%D0%B8%D0%BF%D0%B5%D0%B4%D0%B8%D0%B0:%',
          Row(code='ab', title='%D0%90%D2%B3%D3%99%D1%8B%D0%BD%D2%AD%D2%9B%D0%B0%D1%80%D1%80%D',
          Row(code='ab', title='%D0%91%D1%80%D0%B8%D1%82%D0%B0%D0%BD%D0%B8%D0%B0_%D0%94%D1%83_%',
          Row(code='ab', title='%D0%92%D0%BB%D0%B0%D0%B4%D0%B8%D0%BC%D0%B8%D1%80_%D0%9F%D1%83%']
```

### 1.1.8 Exercise 1.8.

```
In [17]: def f8(rdd):
    # Compute the total number of pageviews for each project
    input_data = rdd.map(lambda line: line.split(" "))
    df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

    # convert to double from string
    df_normal = df_normal.withColumn('size', col('size').cast('double'))
    df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

    # group by project code
    sum = df_normal.groupBy("code").sum("hits").toPandas()

    for index, row in sum.iterrows():
        print("Total number of page views for \"" + str(row[0]) + "\" project = " + str(row[1]))
```

Result of query is commented out because query is very long. Please remove "#" to execute query.

```
In [13]: #f8(rdd)
```

### 1.1.9 Exercise 1.9.

```
In [19]: def f9(rdd):
    # Report the 10 most popular pageviews of all projects , sorted by the total number of pageviews
    input_data = rdd.map(lambda line: line.split(" "))
    df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

    # convert to double from string
    df_normal = df_normal.withColumn('size', col('size').cast('double'))
    df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

    query = df_normal.orderBy(['hits'], ascending=False)

    # clean visualization
    df = query.limit(10).toPandas()
    for index, row in df.iterrows():
        print(row[0], row[1], row[2], row[3])
```

```
In [20]: f9(rdd)
```

```
en.mw en 5466346.0 141180155987.0
es.mw es 695531.0 12261337515.0
ja.mw ja 611443.0 15021588551.0
de.mw de 572119.0 9523069696.0
fr.mw fr 536978.0 11752030020.0
ru.mw ru 466742.0 11847816616.0
it.mw it 400297.0 8176042087.0
```

```
en Main_Page 257915.0 4289970372.0
pt.mw pt 196160.0 4029404403.0
pl.mw pl 176059.0 2782453516.0
```

#### 1.1.10 Exercise 1.10.

```
In [21]: def f10(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

        # Determine the number of page titles that start with the article The.
        # How many of those page titles are not part of the English project?
        query = df_normal.withColumn('The', df_normal['title'].substr(0, 3)=='The')
        query = query.filter(query.The == True)
        print("Number of titles that start with The = "+str(query.count()))
        query = query.filter(query.code!='en')
        print("Number of titles that start with The and are not in \"en\" = "+str(query.c

In [22]: f10(rdd)
```

```
Number of titles that start with The = 48684
Number of titles that start with The and are not in "en" = 11553
```

#### 1.1.11 Exercise 1.11.

```
In [23]: def f11(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

        # convert to double from string
        df_normal = df_normal.withColumn('size', col('size').cast('double'))
        df_normal = df_normal.withColumn('hits', col('hits').cast('double'))

        # Determine the percentage of pages that have only received a single page view in
        query = df_normal.filter(df_normal.hits == 1).count()
        print(str(round(100*query/df_normal.count(),2))+"%")

In [24]: f11(rdd)

79.38%
```

#### 1.1.12 Exercise 1.12.

```
In [29]: def f12(rdd):
        input_data = rdd.map(lambda line: line.split(" "))
        df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])
```

```

# Determine the number of unique terms appearing in the page titles. Note that in
# titles, terms are delimited by "_" instead of a whitespace.
df = df_normal.select(F.split(F.col("title"), "_").rdd.map(lambda r : r[0]))
print(df.flatMap(set).distinct().count())

```

In [30]: f12(rdd)

3491232

### 1.1.13 Exercise 1.13.

```

In [6]: def f13(rdd):
    # Determine the most frequently occurring page title term in this dataset.
    input_data = rdd.map(lambda line: line.split(" "))
    df_normal = input_data.toDF(['code', 'title', 'hits', 'size'])

    df = df_normal.select(F.split(F.col("title"), "_").rdd.map(lambda r : r[0]))
    df = df.flatMap(lambda doc: [(x, 1) for x in doc]).reduceByKey(add)
    df = df.toDF(['word', 'count'])
    x = df.select([max('count')]).collect()[0][0]
    return x

```

In [28]: print(f13(rdd))

215677

## 1.2 Exercise 2

```

In [7]: # set up
sc.stop()
conf = SparkConf().setAppName('appName').setMaster('spark://Jeans-MBP:7077')
sc = SparkContext(conf=conf)

# define rdd
sqlContext = SQLContext(sc)
rdd = sc.textFile(file)

```

### 1.2.1 Small cluster configuration:

```

SPARK_WORKER_CORES=2
SPARK_WORKER_INSTANCES=2
SPARK_WORKER_MEMORY=1g

```

```

In [11]: # Query 3
start_time = time.time()
f3(rdd)

```

```

print("Query 3 takes %s seconds" % (time.time() - start_time))

# Query 5
start_time = time.time()
f5(rdd)
print("Query 5 takes %s seconds" % (time.time() - start_time))

# Query 6
start_time = time.time()
f6(rdd)
print("Query 6 takes %s seconds" % (time.time() - start_time))

# Query 7
start_time = time.time()
f7(rdd)
print("Query 7 takes %s seconds" % (time.time() - start_time))

# Query 13
start_time = time.time()
f13(rdd)
print("Query 13 takes %s seconds" % (time.time() - start_time))

```

```

Query 3 takes 19.41314172744751 seconds
Query 5 takes 44.53948783874512 seconds
Query 6 takes 18.51176691055298 seconds
Query 7 takes 16.548261880874634 seconds
Query 13 takes 69.00603985786438 seconds

```

### 1.2.2 Big cluster configuration:

```

SPARK_WORKER_CORES=8
SPARK_WORKER_INSTANCES=8
SPARK_WORKER_MEMORY=8g

```

```

In [12]: # Query 3
start_time = time.time()
f3(rdd)
print("Query 3 takes %s seconds" % (time.time() - start_time))

# Query 5
start_time = time.time()
f5(rdd)
print("Query 5 takes %s seconds" % (time.time() - start_time))

# Query 6
start_time = time.time()
f6(rdd)

```



```

print("Query 6 takes %s seconds" % (time.time() - start_time))

# Query 7
start_time = time.time()
f7(rdd)
print("Query 7 takes %s seconds" % (time.time() - start_time))

# Query 13
start_time = time.time()
f13(rdd)
print("Query 13 takes %s seconds" % (time.time() - start_time))

```

```

Query 3 takes 18.175821781158447 seconds
Query 5 takes 36.74556493759155 seconds
Query 6 takes 17.101213932037354 seconds
Query 7 takes 15.644676923751831 seconds
Query 13 takes 60.782407999038696 seconds

```

"SPARK\_WORKER\_CORES" is the number of cores per worker.  
 "SPARK\_WORKER\_INSTANCES" is the number of individual workers in the cluster.  
 "SPARK\_WORKER\_MEMORY" is the amount of memory allocated to each cluster.

The big cluster takes significantly more time to compute the queries 3, 5 and 13. This is because as one adds more and more nodes to the cluster, communication cost increases and therefore computing time increases. For queries 6 and 7 the computing time is about the same probably because the queries require less computation power and therefore the extra computation power of the big cluster is not exploited.